

# Bayesian coursework specification for 2021

Data Analytics ECS648U/ ECS784U/ ECS784P

Revised on 25/02/2021 by Dr Anthony Constantinou and Dr Neville Kenneth Kitson.

## 1. Important Dates

- Release date: Thursday **25<sup>th</sup> February** 2021 at 10:00 AM.
- Submission deadline: Wednesday, **28<sup>th</sup> April** 2021 at 10:00 AM.
- Late submission deadline (cumulative penalty applies): Within 7 days after deadline.

### General information:

- i. When submitting coursework online you receive an automated e-mail as proof of submission. Turnitin receipt does not constitute proof of submission. Some students will sometimes upload their coursework and not hit the submit button. Make sure you fully complete the submission process.
- ii. A penalty will be applied automatically by the system for late submissions.
  - a. Your lecturer cannot remove the penalty!
  - b. Penalties can only be challenged via submission of an Extenuating Circumstances (EC) form which can be found on your Student Support page. All the information you need to know is on that page; including how to submit an EC claim along with the deadline dates and full guidelines.
  - c. If you submit an EC form, your case will be reviewed by a panel and the panel will make a decision on the penalty and inform the Module Organiser.
- iii. If you miss both the submission deadline and the late submission deadline, you will automatically receive a score of 0. Extensions can only be granted through approval of an EC claim.
- iv. Submissions via e-mail are not accepted.
- v. It is recommended by the School that we set the deadline at 10:00 AM. Do not wait until the very last moment to submit the coursework.
- vi. Your submission should be a single PDF file.
- vii. For more details on submission regulations, please refer to your relevant handbook.

## 2. Coursework overview

- The coursework is based on the Bayesian material and must be completed individually (group submissions will not be accepted).
- To complete the coursework, follow the tasks below and answer ALL questions enumerated in Section 3. It is recommended that you read the full document *before* you start completing the tasks enumerated below.
- What follows has been tested on Windows and MAC operating systems. There is a compatibility issue with MAC OS (and likely to extend to Linux) which is covered in the Bayesys manual (details below), but which does not influence the coursework submission requirements.

### Task 1: Set up and reading

- a) Visit <http://bayesian-ai.eecs.qmul.ac.uk/bayesys/>
- b) Download the Bayesys user manual.
- c) Set up the project by following the steps in Section 1 of the manual.
- d) Read Section 2 of the manual.
- e) Read Section 3.
- f) Read Section 4.
- g) Skip Section 5.
- h) Read Section 6 and repeat the example.
  - i. MAC and Linux users will not be able to view the PDF graphs shown in Fig 6.1; i.e., the compatibility issue involves the PDF file generator.
  - ii. Skip subsections 6.3, 6.3.1, and 6.4.
- i) Skip Section 7.
- j) Skip Section 8.
- k) Read Section 9.
- l) Skip the appendices.

### Task 2: Determine research area and collate data

You are free to choose or collate your own dataset. You should also determine the dataset size, both in terms of the number of variables and the sample size, relevant to the problem you are analysing. Some areas might require more data than others, and it is up to you to make this decision.

You should address a data-related problem in your professional field or a field you are interested in. If you are motivated in the subject matter the project will be more fun for you and you will likely produce a better report. Section 5 provides a list of data sources you could consider.

You are allowed to reuse the dataset you prepared during the Python coursework, as long as a) your Python coursework submission was *NOT* a group submission, and b) you consider the dataset to be suitable for Bayesian network structure learning (refer to Q1 in Section 3).

Lastly, you are not allowed to reuse datasets from the Bayesys repository for this coursework.

### **Task 3: Prepare your dataset for structure learning**

- a) The Bayesys structure learning system assumes the input data are discrete; e.g., low/medium/high or Yellow/Blue/Green, rather than a continuous range of numbers. If you have a continuous variable in your dataset with integers ranging, for example, from 1 to 100, the algorithm will assume that this variable has 100 different states (and many more if the values are not integer). This will make the dimensionality of the model unmanageable, leading to poor accuracy and high runtime; if this is not clear why, refer to the Conditional Probability Tables (CPTs) in the lecture slides and relevant book material.

You should discretise continuous variables to reduce the number of states to reasonable levels. For example, you could discretise the variable discussed above, with values ranging from 1 to 100, into the five states {"1to20", "21to40", "41to60", "61to80", "81to100"}. If a continuous variable incorporates a small number of different values (e.g., less than 10), it may not need discretisation.

It is up to you to determine whether a variable requires discretisation, as well as the level of discretisation. You are free to follow any approach you wish to discretise the variable, including discretising the variables manually as discussed in the above example. The structure learning accuracy is not expected to be strongly influenced as long as the dimensionality of the data is reasonable with respect to its sample size.

- b) Your dataset must not have missing values (i.e., empty cells). Replace ALL empty cells with the value 'missing' (or use a different relevant name). This forces the algorithm to consider all missing values as an additional state. If missing data follows a pattern, this may or may not help the algorithm to produce a more accurate graph.
- c) Rename your dataset to *trainingData.csv* and place it in folder *Input*.

### Task 3: Draw out your knowledge-based graph

- a) Use your knowledge to produce a knowledge causal graph given the variables in your dataset. You may find it easier if you start drawing the graph by hand.
- b) Record this knowledge in a CSV file following the format of *DAGtrue.csv* as depicted in the Bayesys manual. For an example file, refer to file *DAGtrue\_ASIA.csv* in project directory *Sample input files/Structure learning*.
- c) Rename your knowledge graph file *DAGtrue.csv* and place it in folder *Input*.
- d) Make another copy of the above file, rename it *DAGlearned.csv* and place it in folder *Output*.
- e) Run the Bayesys NetBeans project and make sure your dataset is in folder *Input* and named *trainingData.csv* (as per Task 2c). Under tab *Main*, select *Evaluate graph* and the subprocess *Generate DAGlearned.PDF*. Then hit *Run*.

- i. The system will generate the file *DAGlearned.pdf* in folder *Output*. This is your knowledge graph drawn by the system.

If you are working on MAC/Linux OS, the *DAGlearned.pdf* file is likely to be corrupted. If it is, you can use an online Graphviz editor such as the one available here: <https://edotor.net/>. The Graphviz editor turns a textual representation of a graph into a visual drawing. Use the code shown below, as an example, and edit the code accordingly to be consistent with your *DAGtrue.csv*; e.g., the relationships can be taken directly from the CSV file. The graph should update instantly as you edit the code.

```
digraph {
    Earthquake -> Alarm
    Burglar -> Alarm
    Alarm -> Call
}
```

- ii. This step also generates some information in the terminal window of NetBeans. Save the last three lines as you will need them in answering some of the questions in Section 3; i.e., the line outputs involving Log-Likelihood (LL) score, BIC score and the # of free parameters.

#### Task 4: Run structure learning

- a) Run Bayesys.
  - i. Under tab *Main*, select *Structure learning* with algorithm *SaiyanH* (default selection), select *Evaluate graph* and the subprocess *Generate DAGlearned.PDF*.
  - ii. Under tab *Learning* select *Save associational scores*.
  - iii. Under tab *Evaluation*, make sure that all metrics are selected (they should be selected by default).
  - iv. Under tab *Main*, hit *Run*.
    - i. If your dataset consists of more than 50 variables and/or more than 100k samples (note: there is no requirement for your data to be this big or of a particular size), the learning process may take a while to complete, and this also depends on the number of states per variable (i.e., dimensionality of the data). For smaller datasets, this should complete within a few seconds or minutes. The text output generated in the terminal window of NetBeans indicates the status of learning; i.e., SaiyanH consists of three learning phases and these are reported in the terminal window of NetBeans.
- b) Once the above process completes, the following outputs are generated by the system, which you should save as you will need it in answering some of the questions in Section 3:
  - i. Text output in the terminal window of NetBeans.
  - ii. Files *DAGlearned.csv* and *DAGlearned.pdf* in folder *Output*. Both these files represent your learned graph. As stated Task 3e, you may have to use the online Graphviz editor if you work on MAC/Linux and *DAGlearned.pdf* is corrupted.
  - iii. Four CSV files in directory *Output/SaiyanH*.

c) Repeat the above process for HC and TABU algorithms and save all output information.

- i. Note that subprocess *Save associational scores* does not apply to HC and TABU – so skip this activity.
- ii. Keep in mind that, as stated in the Bayesys manual, Bayesys overwrites files in folder *Output* every time it runs. If you are experimenting and running the algorithms multiple times, you need to remember to rename the files so they are not overwritten by the next run.

Also if you happen to have one of the output files open – for example, looking at *DAGlearned.pdf* in Adobe Reader, and run Bayesys *Evaluate*, then Bayesys will not overwrite it (because Adobe has it locked) so in this case the output file will not reflect the latest run.

### 3. Questions

Important information:

- You should answer ALL questions.
- In your answer sheet, ensure you clearly indicate which answer corresponds to what question. For example:  
*“Answer 2: The steps I followed to produce the knowledge graph are...”*
- DO NOT exceed the maximum number of words indicated for each question. For example, if a question restricts the answer to 100 words, only the first 100 words will be considered in marking the answer.
- Answer the questions in your own words. Copying text from relevant resources will not give you many marks.
- Submission should be a single file containing all your answers. Dataset and other relevant files are not needed.
- Marking is out of 50.
- Refer to Section 4 for the coursework timeline which specifies at which point of this course you are expected to be able to answer each question.

**Question 1:** Discuss the research area and the dataset you have selected or collated for this coursework, along with pointers to your data sources. Screen-capture part of the dataset and present it here as a Figure (e.g., if your dataset contains 15 variables and 1,000 samples, you could show the first 10 columns and a small part of the sample size). Explain why you considered this dataset to be suitable for BN structure learning, and what questions do you expect structure learning to answer.

**Maximum number of words: 200**

**Marks: 5**

**Question 2:** Present the knowledge graph and describe the steps you have followed to produce the graph. For example, what information did you use? Did you refer to the appropriate literature to obtain the necessary knowledge or did you consider your current knowledge to be sufficient for this problem? If you referred to the literature to obtain additional information, provide links to the papers you have read and very briefly describe the knowledge gained from each paper. If you did not refer to the literature, justify why you consider your knowledge to be sufficient.

[NOTE: It is possible to obtain maximum marks without referring to the literature, as long as you clearly justify why this step was not needed].

**Maximum number of words: 250**

**Marks: 5**

**Question 3:** Read the research paper entitled “*Learning Bayesian networks that enable full propagation of evidence*”. This paper describes and evaluates the SaiyanH algorithm. It can be downloaded from this link: <https://ieeexplore.ieee.org/abstract/document/9136714>

Investigate the four CSV files generated in Task 4 during structure learning with SaiyanH. List the number of scores generated in each CSV file. For example, if *marginalDep.csv* has 100 rows of scores, then you should write ‘100’ for that particular file. Explain the different quantities in scores generated in each file. Why do you think there are, for example, 100 scores in the file *marginalDep.csv* and, for example, 200 scores in the file *conditionalInsignificance.csv*?

**Maximum number of words: 100**

**Marks: 5**

**Question 4a:** Refer to the outputs generated in the terminal window of NetBeans during Task 4. Copy, in your answer, the terminal outputs as a Figure, for all the three algorithms. If the information does not fit well into a single figure, you can split it into multiple figures for each algorithm. Make sure the figures show the number of variables and sample size (found under *Training data info* in the terminal), the scores for Precision, Recall, F1, SHD, BSF, LL, BIC, the # of free parameters (found under *Evaluation*), and the elapsed runtime (found under *Structure learning*), for all the three algorithms.

Refer to the F1, SHD and BSF scores produced by SaiyanH and compare them to the related scores shown in Fig 2 of the related research paper (link already provided in Q3). Are your scores mostly lower, on par, or higher (in general) compared to those shown in Fig 2 of the research paper, and with respect to SaiyanH (ignore the results produced by the other algorithms)? Which results did you expect (if any) and which did you not expect (if any)? Explain why.



**Maximum number of words: 200**

**Marks: 5**

**Question 4b:** Further to Q4a, compare the F1, SHD and BSF scores generated across all the three algorithms. Rank the three algorithms by score performance. Which results did you expect (if any) and which did you not expect (if any)? Explain why.

**Maximum number of words: 200**

**Marks: 5**

**Question 5:** Refer to your elapsed structure learning runtime for SaiyanH and compare it to the runtime shown in Table 3 of the related research paper (link already provided in Q3). Indicate whether your results are consistent or not with the results shown in Table 3, and explain why.

**Maximum number of words: 100**

**Marks: 5**

**Question 6:** Compare the BIC scores generated at Task 4, across all the three algorithms, with the BIC score generated at Task 3. What do you understand from the difference between those four scores? Which results did you expect (if any) and which did you not expect (if any)? Explain why.

**Maximum number of words: 200**

**Marks: 5**

**Question 7:** Compare the # of free parameters generated at Task 4, across all three algorithms, with the # of free parameters generated at Task 3. What do you understand from the difference between these four values? Which results did you expect (if any) and which did you not expect (if any)? Explain why.

**Maximum number of words: 200**

**Marks: 5**

**Question 8:** Refer to Week 11 Lecture and Tutorial 2, and select two information fusion methods to apply to the structure learning process of all three algorithms. Each information fusion method should be applied independently to structure learning. It is up to you to decide how much knowledge/information to provide to each information fusion method.

Complete the table below for all nine experiments; i.e., three structure learning runs without information fusion (these can be taken from your previous answers/tasks), three runs based on your first information fusion selection, and another three runs based on your second information fusion selection. To display the table clearly, you may have to move it to a separate page with Landscape orientation with narrow margins.

Explain the differences in scores with and without knowledge. Which results did you expect (if any) and which did you not expect (if any)? Explain why.

	No knowledge						List 1 <sup>st</sup> knowledge constraint here						List 2 <sup>nd</sup> knowledge constraint here								
Algorithm	F1	SHD	BSF	LL	BIC	# free param	runtime	F1	SHD	BSF	LL	BIC	# free param	runtime	F1	SHD	BSF	LL	BIC	# free param	runtime
SaiyanH																					
HC																					
TABU																					

**Maximum number of words: 350**

**Marks: 10**

## 4. Coursework timeline

The table below illustrates at which point of this course you should be able to gain the knowledge needed to answer each of the questions enumerated in Section 3. Some questions are based on material covered over multiple weeks. Do not underestimate the importance of the weekly reading material – especially the research papers. The tutorials are also expected to be particularly helpful in understanding how to use Bayesys to generate the results needed to answer most of the questions.

Question	Week 4 double lecture (Introduction)	Week 5 lecture (Constraint-based learning)	Week 6 lecture (Score-based learning)	Week 6 tutorial (Structure learning using Bayesys)	Week 11 lecture (Evaluation and information fusion)	Week 11 tutorial (Information fusion using Bayesys)	Week 12 lecture (Limitations, interventions, and revision)
1							
2							
3							
4a&b							
5							
6							
7							
8							

## 5. Data sources

Using public data is the most common choice. If you have access to private data, that is also an option, though you will have to be careful about what results you can release. Some sources of publicly available data are listed below (you don't have to use these sources).

- **UK Covid Data**  
<https://coronavirus.data.gov.uk/>  
Official UK COVID data
- **Data.gov**  
<http://data.gov>  
This is the resource for most government-related data.
- **Socrata**  
<http://www.socrata.com/resources/>  
Socrata is a good place to explore government-related data. Furthermore, it provides some visualization tools for exploring data.
- **US Census Bureau**  
<http://www.census.gov/data.html>  
This site provides information about US citizens covering population data, geographic data, and education.
- **UN3ta**  
<https://data.un.org/>  
UN data is an Internet-based data service which brings UN statistical databases.
- **European Union Open Data Portal**  
<http://open-data.europa.eu/en/data/>  
This site provides a lot of data from European Union institutions.
- **Data.gov.uk**  
<http://data.gov.uk/>  
This site of the UK Government includes the British National Bibliography: metadata on all UK books and publications since 1950.
- **The CIA World Factbook**  
<https://www.cia.gov/library/publications/the-world-factbook/>  
This site of the Central Intelligence Agency provides a lot of information on history, population, economy, government, infrastructure, and military of 267 countries.
- **Health Data**  
Healthdata.gov  
<https://www.healthdata.gov/>  
This site provides medical data about epidemiology and population statistics.

- **NHS Health and Social Care Information Centre**  
<http://www.hscic.gov.uk/home>  
 Health datasets from the UK National Health Service.
- **Social Data**  
 Facebook Graph  
<https://developers.facebook.com/docs/graph-api>  
 Facebook provides this API which allows you to query the huge amount of information that users are sharing with the world.
- **Topsy**  
<http://topsy.com/>  
 Topsy provides a searchable database of public tweets going back to 2006 as well as several tools to analyze the conversations.
- **Google Trends**  
<http://www.google.com/trends/explore>  
 Statistics on search volume (as a proportion of total search) for any given term, since 2004.
- **Likebutton**  
<http://likebutton.com/>  
 Mines Facebook's public data--globally and from your own network--to give an overview of what people "Like" at the moment.
- **Amazon Web Services public datasets**  
<http://aws.amazon.com/datasets>  
 The public data sets on Amazon Web Services provide a centralized repository of public data sets. An interesting dataset is the 1000 Genome Project, an attempt to build the most comprehensive database of human genetic information. Also a NASA database of satellite imagery of Earth is available.
- **DBPedia**  
<http://wiki.dbpedia.org>  
 Wikipedia contains millions of pieces of data, structured and unstructured, on every subject. DBPedia is an ambitious project to catalogue and create a public, freely distributable database allowing anyone to analyze this data.
- **Freebase**  
<http://www.freebase.com/>  
 This community database provides information about several topics, with over 45 million entries.
- **Gapminder**  
<http://www.gapminder.org/data/>  
 This site provides data coming from the World Health Organization and World Bank covering economic, medical, and social statistics from around the world.

- **Google Finance**  
<https://www.google.com/finance>  
Forty years' worth of stock market data, updated in real time.
- **National Climatic Data Center**  
<http://www.ncdc.noaa.gov/data-access/quick-links#loc-clim>  
Huge collection of environmental, meteorological, and climate data sets from the US National Climatic Data Center. The world's largest archive of weather data.
- **WeatherBase**  
<http://www.weatherbase.com/>  
This site provides climate averages, forecasts, and current conditions for over 40,000 cities worldwide.
- **Wunderground**  
<http://www.wunderground.com/>  
This site provides climatic data from satellites and weather stations, allowing you to get all information about the temperature, wind, and other climatic measurements.
- **Football datasets**  
<http://www.football-data.co.uk/>  
This site provides historical data for football matches around the world.
- **Pro-Football-Reference**  
<http://www.pro-football-reference.com/>  
This site provides data about football and several other sports.
- **New York Times**  
<http://developer.nytimes.com/docs>  
Searchable, indexed archive of news articles going back to 1851.
- **Google Books Ngrams**  
<http://storage.googleapis.com/books/ngrams/books/datasetv2.html>  
This source searches and analyses the full text of any of the millions of books digitized as part of the Google Books project.
- **Million Song Data Set**  
<http://aws.amazon.com/datasets/6468931156960467>  
Metadata on over a million songs and pieces of music. Part of Amazon Web Services.